

# Algorithms for Data Science

## CSOR W4246

Eleni Drinea  
*Computer Science Department*

Columbia University

Shortest paths in weighted graphs (Bellman-Ford, Floyd-Warshall)

- 1 Shortest paths in graphs with non-negative edge weights (Dijkstra's algorithm)
  - Implementations
  - Graphs with **negative** edge weights: why Dijkstra fails
- 2 Single-source shortest paths (negative edges): Bellman-Ford
  - A DP solution
  - An alternative formulation of Bellman-Ford
- 3 All-pairs shortest paths (negative edges): Floyd-Warshall

# Today

- 1 Shortest paths in graphs with non-negative edge weights (Dijkstra's algorithm)
  - Implementations
  - Graphs with **negative** edge weights: why Dijkstra fails
- 2 Single-source shortest paths (negative edges): Bellman-Ford
  - A DP solution
  - An alternative formulation of Bellman-Ford
- 3 All-pairs shortest paths (negative edges): Floyd-Warshall

# Graphs with **non-negative** weights

## Input

- ▶ a weighted, directed graph  $G = (V, E, w)$ ; function  $w : E \rightarrow \mathbb{R}^+$  assigns non-negative real-valued weights to edges;
- ▶ a source (**origin**) vertex  $s \in V$ .

**Output:** for every vertex  $v \in V$

1. the length of a shortest  $s$ - $v$  path;
2. a shortest  $s$ - $v$  path.

# Dijkstra's algorithm (Input: $G = (V, E, w), s \in V$ )

**Output:** arrays  $dist, prev$  with  $n$  entries such that

1.  $dist(v)$  stores the length of the shortest  $s$ - $v$  path
2.  $prev(v)$  stores the node before  $v$  in the shortest  $s$ - $v$  path

At all times, maintain a set  $S$  of nodes for which the distance from  $s$  has been determined.

- Initially,  $dist(s) = 0, S = \{s\}$ .
- Each time, add to  $S$  the node  $v \in V - S$  that
  1. has an edge from some node in  $S$ ;
  2. minimizes the following quantity among all nodes  $v \in V - S$

$$d(v) = \min_{u \in S: (u,v) \in E} \{dist(u) + w(u, v)\}$$

- Set  $prev(v) = u$ .

# Implementation

**Dijkstra-v1**( $G = (V, E, w), s \in V$ )

  Initialize( $G, s$ )

$S = \{s\}$

**while**  $S \neq V$  **do**

    Select a node  $v \in V - S$  with at least one edge from  $S$  so that

$$d(v) = \min_{u \in S, (u,v) \in E} \{dist[u] + w(u, v)\}$$

$S = S \cup \{v\}$

$dist[v] = d(v)$

$prev[v] = u$

**end while**

Initialize( $G, s$ )

**for**  $v \in V$  **do**

$dist[v] = \infty$

$prev[v] = NIL$

**end for**

$dist[s] = 0$

# Improved implementation (I)

Idea: Keep a **conservative overestimate** of the true length of the shortest  $s$ - $v$  path in  $dist[v]$  as follows: when  $u$  is added to  $S$ , **update**  $dist[v]$  for all  $v$  with  $(u, v) \in E$ .

Dijkstra-v2( $G = (V, E, w), s \in V$ )

  Initialize( $G, s$ )

$S = \emptyset$

**while**  $S \neq V$  **do**

    Pick  $u$  so that  $dist[u]$  is minimum among all nodes in  $V - S$

$S = S \cup \{u\}$

**for**  $(u, v) \in E$  **do**

      Update( $u, v$ )

**end for**

**end while**

Update( $u, v$ )

**if**  $dist[v] > dist[u] + w(u, v)$  **then**

$dist[v] = dist[u] + w(u, v)$

$prev[v] = u$

**end if**

# Improved implementation (II): binary min-heap

Idea: Use a **priority queue implemented as a binary min-heap**: store vertex  $u$  with key  $dist[u]$ . Required operations: **Insert**, **ExtractMin**; **DecreaseKey** for **Update**; each takes  $O(\log n)$  time.

Dijkstra-v3( $G = (V, E, w), s \in V$ )

  Initialize( $G, s$ )

$Q = \{V; dist\}$

$S = \emptyset$

**while**  $Q \neq \emptyset$  **do**

$u = \text{ExtractMin}(Q)$

$S = S \cup \{u\}$

**for**  $(u, v) \in E$  **do**

      Update( $u, v$ )

**end for**

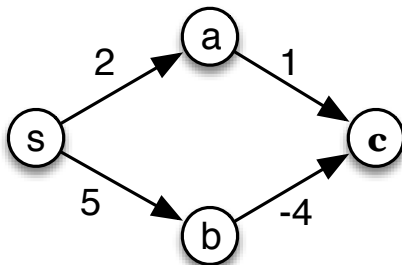
**end while**

**Running time:**  $O(n \log n + m \log n) = O(m \log n)$

*When is Dijkstra-v3() better than Dijkstra-v2()?*

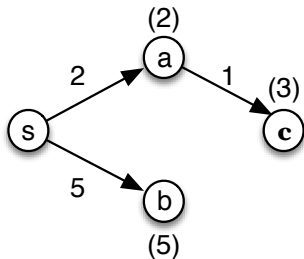


## Example graph with **negative** edge weights

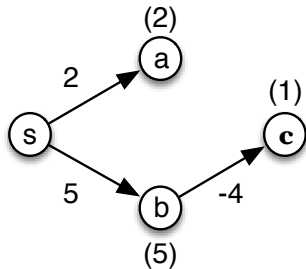


# Dijkstra's output and correct output for example graph

Dijkstra's output



Correct shortest paths




Dijkstra's algorithm will first include  $a$  to  $S$  and then  $c$ , thus missing the shorter path from  $s$  to  $b$  to  $c$ .

# Negative edge weights: why Dijkstra's algorithm fails

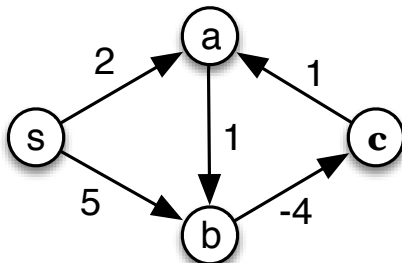
- ▶ Intuitively, a path may start on long edges but then compensate along the way with short edges.
- ▶ Formally, in the proof of correctness of the algorithm, the last statement about  $P$  does not hold anymore: even if the length of path  $P_v$  is smaller than the length of the subpath  $s \rightarrow x \rightarrow y$ , negative edges on the subpath  $y \rightarrow v$  may now result in  $P$  being *shorter* than  $P_v$ .

## Bigger problems in graphs with negative edges?



$dist(a) = ?$

## Bigger problems in graphs with negative edges?



1.  $dist(v)$  goes to  $-\infty$  for every  $v$  on the cycle  $(a, b, c, a)$
  2. **no** solution to shortest paths when negative cycles
- $\Rightarrow$  need to **detect** negative cycles

# Today

- 1 Shortest paths in graphs with non-negative edge weights (Dijkstra's algorithm)
  - Implementations
  - Graphs with **negative** edge weights: why Dijkstra fails
- 2 Single-source shortest paths (negative edges): Bellman-Ford
  - A DP solution
  - An alternative formulation of Bellman-Ford
- 3 All-pairs shortest paths (negative edges): Floyd-Warshall

# Single-source shortest paths, negative edge weights

**Input:** weighted directed graph  $G = (V, E, w)$  with  $w : E \rightarrow \mathbb{R}$ ;  
a source (**origin**) vertex  $s \in V$ .

**Output:**

1. If  $G$  has a negative cycle reachable from  $s$ , answer  
    **“negative cycle in  $G$ ”**.
2. Else, compute for every  $v \in V$ 
  - 2.1 the length of a shortest  $s$ - $v$  path;
  - 2.2 a shortest  $s$ - $v$  path.

# Properties of shortest paths

Suppose the problem **has** a solution for an input graph.

- ▶ *Can there be negative cycles in the graph?*
- ▶ *Can there be positive cycles in the graph?*
- ▶ *Can the shortest paths contain positive cycles?*
- ▶ *Consider a shortest  $s$ - $t$  path; are its subpaths shortest? In other words, does the problem exhibit optimal substructure?*



# Towards a DP solution

**Key observation:** if there are no negative cycles, a path cannot become shorter by traversing a cycle.

## Fact 1.


*If  $G$  has no negative cycles, then there is a shortest  $s$ - $v$  path that is simple, thus has at most  $n - 1$  edges.*

## Fact 2.

*The shortest paths problem exhibits optimal substructure.*

Facts 1 and 2 suggest a DP solution.

# Subproblems



Let

$OPT(i, v)$  = cost of a shortest  $s$ - $v$  path with *at most*  $i$  edges

Consider a shortest  $s$ - $v$  path using at most  $i$  edges.

- ▶ If the path uses at most  $i - 1$  edges, then

$$OPT(i, v) = OPT(i - 1, v).$$

- ▶ If the path uses  $i$  edges, then

$$OPT(i, v) = \min_{x:(x,v) \in E} \{OPT(i - 1, x) + w(x, v)\}.$$

# Recurrence

Let


$OPT(i, v)$  = cost of a shortest  $s$ - $v$  path using *at most*  $i$  edges

Then

$$OPT(i, v) = \begin{cases} 0 & , \text{ if } i = 0, v = s \\ \infty & , \text{ if } i = 0, v \neq s \\ \min \left\{ \begin{array}{l} OPT(i-1, v) \\ \min_{x:(x,v) \in E} \{OPT(i-1, x) + w(x, v)\} \end{array} \right\} & , \text{ if } i > 0 \end{cases}$$

## Example of Bellman-Ford

Compute shortest  $s$ - $v$  paths in the graph below, for all  $v \in V$ .



# Pseudocode

$n \times n$  dynamic programming table  $M$  such that  
 $M[i, v] = OPT(i, v)$ .

**Bellman-Ford**( $G = (V, E, w), s \in V$ )

**for**  $v \in V$  **do**

$M[0, v] = \infty$

**end for**

$M[0, s] = 0$

**for**  $i = 1, \dots, n - 1$  **do**

**for**  $v \in V$  (*in any order*) **do**

$$M[i, v] = \min \left\{ \begin{array}{l} M[i - 1, v] \\ \min_{x: (x, v) \in E} \left\{ M[i - 1, x] + w(x, v) \right\} \end{array} \right\}$$

**end for**

**end for**

# Running time & Space

- ▶ **Running time:**  $O(nm)$
- ▶ **Space:**  $\Theta(n^2)$  —can be improved (*coming up*)
- ▶ To reconstruct actual shortest paths, also keep array *prev* of size  $n$  such that

$prev[v]$  = predecessor of  $v$  in current shortest  $s$ - $v$  path.

# Improving space requirements

Only need two rows of  $M$  at all times.

△ Actually, only need one (see Remark 1)! Thus drop the index  $i$  from  $M[i, v]$  and only use it as a counter for #repetitions.

$$M[v] = \min \left\{ M[v], \min_{x:(x,v) \in E} \{ M[x] + w(x, v) \} \right\}$$


## Remark 1.

*Throughout the algorithm,  $M[v]$  is the length of some  $s$ - $v$  path. After  $i$  repetitions,  $M[v]$  is no larger than the length of the current shortest  $s$ - $v$  path with at most  $i$  edges.*

**Early termination condition:** if at some iteration  $i$  **no** value in  $M$  changed, then stop (*why?*)

\* This allows us to detect negative cycles!

# An alternative way to view Bellman-Ford



- ▶ Let  $P = (s = v_0, v_1, v_2, \dots, v_k = v)$  be a shortest  $s$ - $v$  path.
- ▶ Then  $P$  can contain at most  $n - 1$  edges.
- ▶ *How can we correctly compute  $\text{dist}(v)$  on this path?*



## Key observations about subroutine `Update`( $u, v$ )

Recall subroutine `Update` from Dijkstra's algorithm:

$$\text{Update}(u, v) : \text{dist}(v) = \min\{\text{dist}(v), \text{dist}(u) + w(u, v)\}$$

### Fact 3.

*Suppose  $u$  is the last node before  $v$  on the shortest  $s$ - $v$  path, and suppose  $\text{dist}(u)$  has been correctly set. The call `Update`( $u, v$ ) returns the correct value for  $\text{dist}(v)$ .*

### Fact 4.

*No matter how many times `Update`( $u, v$ ) is performed, it will never make  $\text{dist}(v)$  too small. That is, `Update` is a **safe** operation: performing few extra updates can't hurt.*

# Performing the correct sequence of updates

Suppose we update the edges on the shortest path  $P$  **in the order they appear on the path** (though not necessarily consecutively). Hence we update

$$(s, v_1), (v_1, v_2), (v_2, v_3), \dots, (v_{k-1}, v).$$

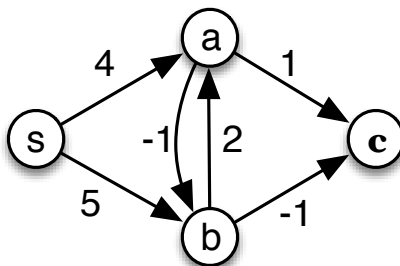
This sequence of updates correctly computes  $\text{dist}(v_1)$ ,  $\text{dist}(v_2)$ ,  $\dots$ ,  $\text{dist}(v)$  (by induction and Fact 3).

*How can we guarantee that this specific sequence of updates occurs?*

## A concrete example

Consider the shortest  $s$ - $b$  path, which uses edges  $(s, a), (a, b)$ .

*How can we guarantee that our algorithm will update these two edges **in this order**? (More updates in between are allowed.)*



# Bellman-Ford algorithm

Update all  $m$  edges in the graph,  $n - 1$  times in a row!

- ▶ By Fact 4, it is ok to update an edge several times in between.
- ▶ All we need is to update the edges on the path in this **particular order**. This is guaranteed if we update all edges  $n - 1$  times in a row.

# Pseudocode

We will use **Initialize** and **Update** from Dijkstra's algorithm.

**Initialize**( $G, s$ )

**for**  $v \in V$  **do**

$dist[v] = \infty$

$prev[v] = NIL$

**end for**

$dist[s] = 0$

**Update**( $u, v$ )

**if**  $dist[v] > dist[u] + w(u, v)$  **then**

$dist[v] = dist[u] + w(u, v)$

$prev[v] = u$


**end if**

# Bellman-Ford

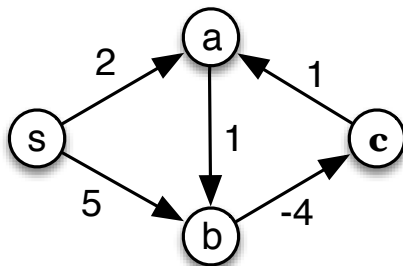
```
Bellman-Ford( $G = (V, E, w), s$ )  
  Initialize( $G, s$ )  
  for  $i = 1, \dots, n - 1$  do  
    for  $(u, v) \in E$  do  
      Update( $u, v$ )  
    end for  
  end for
```

Running time? Space?

## Detecting negative cycles



# Detecting negative cycles



1.  $dist(v)$  goes to  $-\infty$  for every  $v$  on the cycle.
2. Any shortest  $s$ - $v$  path can have at most  $n - 1$  edges.
3. Update all edges  $n$  times (instead of  $n - 1$ ): if  $dist(v)$  changes **for any**  $v \in V$ , then there is a negative cycle.



# Today

- 1 Shortest paths in graphs with non-negative edge weights (Dijkstra's algorithm)
  - Implementations
  - Graphs with **negative** edge weights: why Dijkstra fails
- 2 Single-source shortest paths (negative edges): Bellman-Ford
  - A DP solution
  - An alternative formulation of Bellman-Ford
- 3 All-pairs shortest paths (negative edges): Floyd-Warshall

# All pairs shortest-paths

- ▶ **Input:** a directed, weighted graph  $G = (V, E, w)$  with real edge weights
- ▶ **Output:** an  $n \times n$  matrix  $D$  such that

$$D[i, j] = \text{length of shortest path from } i \text{ to } j$$


# Solving all pairs shortest-paths

1. Straightforward solution: run Bellman–Ford once for every vertex ( $O(n^2m)$  time).
2. Improved solution: Floyd–Warshall’s dynamic programming algorithm ( $O(n^3)$  time).

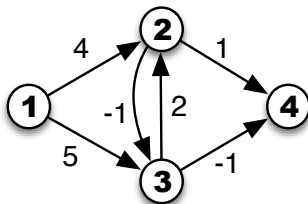
# Towards a DP formulation

- ▶ Consider a shortest  $s$ - $t$  path  $P$ .
- ▶ This path uses some **intermediate** vertices: that is, if  $P = (s, v_1, v_2, \dots, v_k, t)$ , then  $v_1, \dots, v_k$  are intermediate vertices.
- ▶ For simplicity, relabel the vertices in  $V$  as  $\{1, 2, 3, \dots, n\}$  and consider a shortest  $i$ - $j$  path where **intermediate vertices may only be from  $\{1, 2, \dots, k\}$** .
- ▶ **Goal:** compute the length of a shortest  $i$ - $j$  path for every pair of vertices  $(i, j)$ , using  $\{1, 2, \dots, n\}$  as intermediate vertices.

# Example




Rename {s, a, b, c} as **{1, 2, 3, 4}**



# Examples of shortest paths

Shortest **(1, 2)**-path using  $\{\}$  or  $\{1\}$  is **P**.


Shortest **(1, 2)**-path using  $\{1,2,3,4\}$  is **P**.



Shortest **(1, 3)**-path using  $\{\}$  or  $\{1\}$  is **P'**.

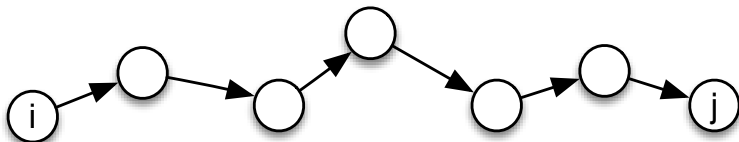
Shortest **(1, 3)**-path using  $\{1,2\}$  or  $\{1,2,3\}$  is **P<sub>1</sub>**.

Shortest **(1, 3)**-path using  $\{1,2,3,4\}$  is **P<sub>1</sub>**.



## A shortest $i$ - $j$ path using nodes from $\{1, \dots, k\}$

Consider a shortest  $i$ - $j$  path  $P$  where intermediate nodes may only be from the set of nodes  $\{1, 2, \dots, k\}$ .

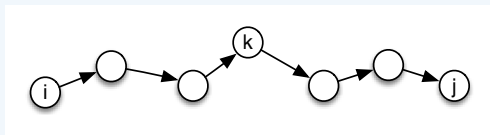


**Fact:** any subpath of  $P$  must be shortest itself.

## A useful observation

Focus on the last node  $k$  from the set  $\{1, 2, \dots, k\}$ . Either

1.  $P$  completely avoids  $k$ : then a shortest  $i$ - $j$  path with intermediate nodes from  $\{1, \dots, k\}$  is the same as a shortest  $i$ - $j$  path with intermediate nodes from  $\{1, \dots, k - 1\}$ .
2. Or,  $k$  is an intermediate node of  $P$ .



Decompose  $P$  into an  $i$ - $k$  subpath  $P_1$  and a  $k$ - $j$  subpath  $P_2$ .

- $P_1, P_2$  are shortest subpaths themselves.
- All intermediate nodes of  $P_1, P_2$  are from  $\{1, \dots, k - 1\}$ .



Let

$$OPT_k(i, j) = \text{cost of shortest } i - j \text{ path } P \text{ using} \\ \{1, \dots, k\} \text{ as intermediate vertices}$$

1. Either  $k$  does not appear in  $P$ , hence

$$OPT_k(i, j) = OPT_{k-1}(i, j)$$

2. Or,  $k$  appears in  $P$ , hence

$$OPT_k(i, j) = OPT_{k-1}(i, k) + OPT_{k-1}(k, j)$$

Hence

$$OPT_k(i, j) = \begin{cases} w(i, j) & , \text{ if } k = 0 \\ \min \begin{cases} OPT_{k-1}(i, j) \\ OPT_{k-1}(i, k) + OPT_{k-1}(k, j) \end{cases} & , \text{ if } k \geq 1 \end{cases}$$

We want  $OPT_n(i, j)$ .

Time/space requirements?

# Floyd-Warshall on example graph

Let  $D_k[i, j] = OPT_k(i, j)$ .

 $D_0 =$ 

0	4	5	$\infty$
$\infty$	0	-1	1
$\infty$	2	0	-1
$\infty$	$\infty$	$\infty$	0

 $D_1 =$ 

0	4	5	$\infty$
$\infty$	0	-1	1
$\infty$	2	0	-1
$\infty$	$\infty$	$\infty$	0

 $D_2 =$ 

0	4	3	5
$\infty$	0	-1	1
$\infty$	2	0	-1
$\infty$	$\infty$	$\infty$	0

 $D_3 =$ 

0	4	3	2
$\infty$	0	-1	-2
$\infty$	2	0	-1
$\infty$	$\infty$	$\infty$	0

# Space requirements

- ▶ A single  $n \times n$  dynamic programming table  $D$ , initialized to  $w(i, j)$  (the adjacency matrix of  $G$ ).
- ▶ Let  $\{1, \dots, k\}$  be the set of intermediate nodes that may be used for the shortest  $i$ - $j$  path.
- ▶ After the  $k$ -th iteration,  $D[i, j]$  contains the length of some  $i$ - $j$  path that is no larger than the length of the shortest  $i$ - $j$  path using  $\{1, \dots, k\}$  as intermediate nodes.

# The Floyd-Warshall algorithm

```
Floyd-Warshall( $G = (V, E, w)$ )  
  for  $k = 1$  to  $n$  do  
    for  $i = 1$  to  $n$  do  
      for  $j = 1$  to  $n$  do  
         $D[i, j] = \min(D[i, j], D[i, k] + D[k, j])$   
      end for  
    end for  
  end for
```

- ▶ Running time:  $O(n^3)$
- ▶ Space:  $\Theta(n^2)$