

Algorithms for Data Science

CSOR W4246

Eleni Drinea
Computer Science Department

Columbia University

Hashing, bloom filters

Outline

- 1 Hashing
- 2 Analyzing hash tables using balls and bins
- 3 Saving space: hashing-based fingerprints
- 4 Bloom filters

Today

- 1 Hashing
- 2 Analyzing hash tables using balls and bins
- 3 Saving space: hashing-based fingerprints
- 4 Bloom filters

The problem

A data structure maintaining a dynamic subset S of a huge universe U .

- ▶ Typically, $|S| \ll |U|$

The data structure should support

- ▶ efficient **insertion**
- ▶ efficient **deletion**
- ▶ efficient **search**

We will call such a data structure a **dictionary**.

Dictionary data structure

A dictionary maintains a subset S of a universe U so that inserting, deleting and searching is efficient.

Operations supported by a dictionary

1. **Create()**: initialize a dictionary with $S = \emptyset$
2. **Insert(x)**: add x to S , if $x \notin S$
 - ▶ additional information about x might be stored in the dictionary as part of a record for x
3. **Delete(x)**: delete x from S , if $x \in S$
4. **Lookup(x)**: determine if $x \in S$

A concrete example

We want to maintain a dynamic list of 250 IP addresses

- ▶ e.g., these correspond to addresses of currently active customers of a Web service
- ▶ each IP address consists of 32 bits, e.g. 128.32.168.80

The challenge: U is enormous, that is, $|U| \gg |S|$

1. Maintain **array** S of size $|U|$ such that $S[i] = 1$ if and only if $i \in S$

- ▶ Insert, Delete, Lookup require $O(1)$ time

Can't store an array of size anywhere close to $|U|$!

- ▶ S should have $|U| = 2^{32} \approx 4$ billion entries
- ▶ S would be mostly empty (huge waste of space)

2. Store S in a **linked list**

- ▶ Space: proportional to $|S| = 250$
- ▶ Time for Lookup: proportional to $|S|$; **too slow**

Can we support fast Insert, Delete, Lookup (as in array implementation) but only use space proportional to $|S|$ (linked list implementation)?

Work with array of size $|S|$ rather than one of size $|U|$

Idea: assign a short *nickname* to each element in U

- ▶ Each of the 2^{32} IP addresses is assigned a number between 1 and $|S| = 250$
 - ▶ range will be slightly adjusted
- ▶ Total amount of storage: approximately $|S|$, **independent of $|U|$**
- ▶ If not too many IP addresses per nickname, then Lookup is **efficient** (*details coming up*)

How can we assign a short name?

By **hashing**: use a hash function $h : U \rightarrow \{0, \dots, n - 1\}$

- ▶ Typically, $n \ll |U|$ and is close to $|S|$

For example,

- ▶ $h : \{0, \dots, 2^{32} - 1\} \rightarrow \{0, \dots, 249\}$
- ▶ IP address x gets name $h(x)$
- ▶ Hash table H of size 250: store address x at entry $h(x)$

So **Insert**(x) takes constant time. *What if we try to insert $y \neq x$, with $h(x) = h(y)$?*


Collision: elements $x \neq y$ such that $h(x) = h(y)$

Easiest way to deal with collisions: **chain hashing**

- ▶ Entry i in the hash table is a **linked list** of elements x such that $h(x) = i$
- ▶ Alternatively, can think of every entry in the hash table as a **bin** containing the elements that hash to the same location

Chain hashing

Maintain a linked list at $H[i]$ for all x such that $h(x) = i$.



Chain hashing: running time for $\text{Lookup}(x)$

Time for $\text{Lookup}(x)$:

1. time to compute $h(x)$; typically, constant
2. time to scan the linked list at position $h(x)$ in hash table
 - ▶ proportional to the *length* of the linked list at $h(x)$, which is proportional to the # elements that collide with x

Goal: find a hash function that “spreads out” the elements well

Simple hash functions might not work

Consider the following two simple hash functions that hash an IP address x from $\{0, \dots, 2^{32} - 1\}$ to $\{0, \dots, 255\}$:

- ▶ assign the last 8 bits of x as its name
- ▶ assign the first 8 bits of x as its name

Remark 1.

*Nothing is **inherently** wrong with these hash functions: the problem is that our 250 IP addresses might not be drawn uniformly at random from among all 2^{32} possibilities.*

No single hash function can work well on *all* data sets

- ▶ **Fix** the hash function h .
- ▶ h distributes $|U|$ elements into n names.
- ⇒ exists data set of at least $\frac{|U|}{n}$ elements that all map to the same name
- ⇒ if our customers come from this data set, lots of collisions

Fact: for any **fixed** (**deterministic**) $h : U \rightarrow \{0, 1, \dots, n - 1\}$ where $|U| \geq n^2$, there exists some set S of n elements that all map to the same position.

Randomization can help

- ▶ **Extreme example:** for every $0 \leq j \leq n - 1$, assign name j to element x with probability $\frac{1}{n}$.
 - ▶ Fix $x, y \in U$. Then $\Pr[h(x) = h(y)] = \frac{1}{n}$.
 - ▶ **This doesn't quite work.** (Think $\text{Lookup}(x)$: *where is x ?*)
 - ▶ However, intuitively, hash functions that spread things around in a *random* way can effectively reduce collisions.
- ⇒ Trade-off in hash function design: h must be “random” to scatter things around for all inputs but still be a function

Goal: design h that allows for efficient dictionary operations
with high probability

A careful use of randomization

- ▶ Randomize over the **choice** of the hash function from a suitable **class of functions** into $[0, n - 1]$ (*details coming up*)
- ▶ h must have a **compact** representation

Universal hash function

Idea: choose h **at random** from a carefully selected class of functions H with the following properties:

1. h behaves almost like a completely random hash function.
 - ▶ For $x, y \in U$. The probability that a randomly chosen $h \in H$ satisfies $h(x) = h(y)$ is at most $1/n$.
2. Can select a random h efficiently.
3. Given h , can compute $h(x)$ efficiently.

Such hash functions are called **universal**; their design relies on number theoretic facts.

Example of universal hash function

- ▶ Pick a prime p close to $|S| = 250$; set $n = p$
 - ▶ E.g., pick $p = 257$; set the size n of the hash table to 257
- ▶ Look at IP address x as (x_1, x_2, x_3, x_4) , where x_1, x_2, x_3, x_4 are integers mod n .
- ▶ Define $h : U \rightarrow \{0, 1, \dots, n - 1\}$ as follows:
 - ▶ Choose a_1, a_2, a_3, a_4 randomly from $\{0, 1, \dots, n - 1\}$
 - ▶ E.g., $a_1 = 80, a_2 = 35, a_3 = 168, a_4 = 220$
 - ▶ Map IP address x to $h(x) = \left(\sum_{i=1}^4 a_i x_i \right) \bmod n$
 - ▶ E.g., $x = 128.32.168.80$,
 $h(x) = (80 \cdot 128 + 35 \cdot 32 + 168 \cdot 168 + 220 \cdot 80) \bmod 257$

h is a universal hash function

Claim 1.

*Consider any pair $x = (x_1, x_2, x_3, x_4)$, $y = (y_1, y_2, y_3, y_4)$.
If a_1, \dots, a_4 are chosen uniformly at random from
 $\{0, \dots, n-1\}$, then*

$$\Pr[h_a(x_1, \dots, x_4) = h_a(y_1, \dots, y_4)] = \frac{1}{n}$$

The proof relies on elementary number theory.

Corollary 1.

*Fix $x \in U$. The expected #elements colliding with x is less than
1. Hence the expected lookup time is constant.*

Ideal hash functions

From now on, *assume a completely random hash function exists.*

△ Does not exist! But can provide a good rough idea of how hashing schemes perform in practice.

- ▶ Let $h : U \rightarrow \{0, 1, \dots, n - 1\}$ be a completely random (ideal) hash function. For all $x \in U$, $0 \leq j \leq n - 1$

$$\Pr[h(x) = j] = \frac{1}{n}$$

Remark 2.

$h(x)$ is **fixed** for every x : it just takes **one** of the n possible values with equal probability.

Today

- 1 Hashing
- 2 Analyzing hash tables using balls and bins
- 3 Saving space: hashing-based fingerprints
- 4 Bloom filters

Hashing modeled as a balls and bins problem

Q1: How many elements can we insert in the hash table before it is more likely than not that there is a collision?

Hashing modeled as a balls and bins problem

Q1: How many elements can we insert in the hash table before it is more likely than not that there is a collision?

This is just an **occupancy problem**!

Hashing modeled as a balls and bins problem

Q1: How many elements can we insert in the hash table before it is more likely than not that there is a collision?

Occupancy problems, revisited: find the distribution of balls into bins when m balls are thrown independently and uniformly at random into n bins.

Hashing modeled as a balls and bins problem

Q1: How many elements can we insert in the hash table before it is more likely than not that there is a collision?

Occupancy problems, revisited: find the distribution of balls into bins when m balls are thrown independently and uniformly at random into n bins.

Hashing as an occupancy problem:

- ▶ balls correspond to elements from U
- ▶ bins are slots in the hash table
- ▶ each ball falls into one of the n bins independently and with probability $1/n$

Hashing modeled as a balls and bins problem

Q1: How many elements can we insert in the hash table before it is more likely than not that there is a collision?

Hashing as an occupancy problem:

- ▶ balls correspond to elements from U
- ▶ bins are slots in the hash table
- ▶ each ball falls into one of the n bins independently and with probability $1/n$

Q1 (rephrased): How many balls can we throw before it is more likely than not that some bin contains at least two balls?

Answer: $\Omega(\sqrt{n})$ (see the birthday paradox)

Towards analyzing time/space efficiency of hash table

- ▶ *What is the expected time for $\text{Lookup}(x)$?*
- ▶ *What is the expected wasted space in the hash table?*
- ▶ *What is the worst-case time for $\text{Lookup}(x)$?*

Towards analyzing time/space efficiency of hash table

- ▶ *What is the expected time for $\text{Lookup}(x)$?*
Corresponds to expected load of a bin.
- ▶ *What is the expected wasted space in the hash table?*
Corresponds to expected number of empty bins.
- ▶ *What is the worst-case time for $\text{Lookup}(x)$?*
Corresponds to load of the fullest bin.

Towards analyzing time/space efficiency of hash table

For $n = m$

- ▶ *What is the expected time for $\text{Lookup}(x)$? $O(1)$.*
- ▶ *What is the expected wasted space in the hash table?
At least a third of the slots are empty.*
- ▶ *What is the worst-case time for $\text{Lookup}(x)$, with high probability?
 $\Theta(\ln n / \ln \ln n)$, with high probability.*

Max load in any bin, with high probability (case $m = n$)

Proposition 1.

When throwing n balls into n bins uniformly and independently at random, the maximum load in any bin is $\Theta(\ln n / \ln \ln n)$ with probability close to 1 as n grows large.

Two-sentence sketch of the proof.

1. Upper bound the probability that **any** bin contains more than k balls by a union bound:
$$\sum_{j=1}^n \sum_{\ell=k}^n \binom{n}{\ell} \left(\frac{1}{n}\right)^\ell \left(1 - \frac{1}{n}\right)^{n-\ell}.$$
2. Compute the smallest possible k^* such that the probability above is less than $1/n$ (which becomes negligible as n grows large).



Today

- 1 Hashing
- 2 Analyzing hash tables using balls and bins
- 3 Saving space: hashing-based fingerprints**
- 4 Bloom filters

A password checker

- ▶ We want to maintain a dictionary for a set S of 2^{16} **bad** passwords so that, when a user tries to set up a password, we can check as quickly as possible if it belongs to S and reject it.
- ▶ We assume that each password consists of 8 ASCII characters
 - ▶ hence each password requires 8 bytes (64 bits) to represent

A dictionary data structure that uses less space

Let S be the set of **bad** passwords.

Input: a 64-bit password x , and a query of the form
*“Is x a **bad** password?”*

Output: a dictionary data structure for S that answers queries as above and

- ▶ is **small**: uses **less space** than explicitly storing all bad passwords
- ▶ allows for erroneous **yes** answers occasionally
 - ▶ that is, we occasionally answer “ $x \in S$ ” even though $x \notin S$

Approximate set membership

The password checker belongs to a broad class of problems, called *approximate set membership* problems.

Input: a large set $S = \{s_1, \dots, s_m\}$, and queries of the form “Is $x \in S$?”

We want a dictionary for S that is **small** (smaller than the explicit representation provided by a hash table).

To achieve this, we allow for some probability of error

- ▶ **False positives:** answer **yes** when $x \notin S$
- ▶ **False negatives:** answer **no** when $x \in S$

Output: small probability of false positives, no false negatives

Fingerprints: hashing for saving space

- ▶ Use a hash function $h : \{0, \dots, 2^{64} - 1\} \rightarrow \{0, \dots, 2^{32} - 1\}$ to map each password into a 32 bit string.
- ▶ This string will serve as a short *fingerprint* of the password.
- ▶ Keep the *fingerprints* in a sorted list.
- ▶ To check if a proposed password is **bad**:
 1. calculate its *fingerprint*
 2. binary search for the *fingerprint* in the list of fingerprints; if found, declare the password **bad** and ask the user to enter a new one.

Setting the length b of the fingerprint

Why did we map passwords to 32-bit fingerprints?

Motivation: make fingerprints long enough so that the false positive probability is acceptable

Let b be the number of bits used by our hash function to map the m bad passwords into fingerprints, thus

$$h : \{0, 1, \dots, 2^{64} - 1\} \rightarrow \{0, \dots, 2^b - 1\}$$

We will choose b so that the probability of a false positive is acceptable, e.g., at most $1/m$.

Determining the false positive probability

There are 2^b possible strings of length b .

Let x be a **good** password.

Fix a $y \in S$ (recall that all m passwords in S are **bad**).

- ▶ $\Pr[x \text{ has the same fingerprint as } y] = 1/2^b$
- ▶ $\Pr[x \text{ does not have the same fingerprint as } y] = 1 - 1/2^b$
- ▶ let $p = 1 - 1/2^b$
- ▶ $\Pr[x \text{ does not have the same fingerprint as any } w \in S] = p^m$
- ▶ $\Pr[x \text{ has the same fingerprint as some } w \in S] = 1 - p^m$

Hence the false positive probability is

$$1 - p^m = 1 - (1 - 1/2^b)^m \approx 1 - e^{-m/2^b}$$

Constant false positive probability and bound for b

To make the probability of a false positive less than, say, a constant c , we require

$$1 - e^{-m/2^b} \leq c \Rightarrow b \geq \log_2 \frac{m}{\ln(1/(1-c))}.$$

So $b = \Omega(\log_2 \frac{m}{\ln(1/(1-c))})$ bits.

Improved false positive probability and bound for b

Now suppose we use $b = 2 \log_2 m$.

Plugging back into the original formula for the probability of false positive, which is $1 - (1 - 1/2^b)^m$, we get

$$1 - \left(1 - \frac{1}{m^2}\right)^m \leq 1 - \left(1 - \frac{1}{m}\right) = \frac{1}{m}$$

Thus if our dictionary has $|S| = m = 2^{16}$ bad passwords, using a hash function that maps each of the m passwords to 32 bits yields a false positive probability of about $1/2^{16}$.

Today

- 1 Hashing
- 2 Analyzing hash tables using balls and bins
- 3 Saving space: hashing-based fingerprints
- 4 Bloom filters**

Fast approximate set membership

Input: a *large* set S , and queries of the form “ $Is\ x \in S?$ ”

Fast approximate set membership

Input: a *large* set S , and queries of the form “ $Is\ x \in S?$ ”

We want a **data structure** that answers the queries

- ▶ **fast** (faster than searching in S)
- ▶ is **small** (smaller than the explicit representation provided by hash table)

Fast approximate set membership

Input: a *large* set S , and queries of the form “Is $x \in S$?”

We want a **data structure** that answers the queries

- ▶ **fast** (faster than searching in S)
- ▶ is **small** (smaller than the explicit representation provided by hash table)

To achieve the above, allow for some probability of error

- ▶ **False positives:** answer **yes** when $x \notin S$
- ▶ **False negatives:** answer **no** when $x \in S$

Fast approximate set membership

Input: a *large* set S , and queries of the form “Is $x \in S$?”

We want a **data structure** that answers the queries

- ▶ **fast** (faster than searching in S)
- ▶ is **small** (smaller than the explicit representation provided by hash table)

To achieve the above, allow for some probability of error

- ▶ **False positives:** answer **yes** when $x \notin S$
- ▶ **False negatives:** answer **no** when $x \in S$

Output: small probability of false positives, no false negatives

Bloom filter

A Bloom filter consists of:

1. an array B of n **bits**, initially all set to 0.

$B =$

| | | | | | | | | | | | | | | | | |
|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|
| 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|

2. k independent random hash functions h_1, \dots, h_k with range $\{0, 1, \dots, n - 1\}$.

A basic Bloom filter supports

- ▶ $\text{Insert}(x)$
- ▶ $\text{Lookup}(x)$

Representing a set $S = \{x_1, \dots, x_m\}$ using a Bloom filter

SetupBloomFilter(S, h_1, \dots, h_k)

Initialize array B of size n to all zeros

for $i = 1$ to m **do**

 Insert(x_i)

end for

Insert(x)

for $i = 1$ to k **do**

 compute $h_i(x)$




 set $B[h_i(x)] = 1$

end for

Remark: an entry of B may be set multiple times; only the first change has an effect.

Setting up the Bloom filter

$S = \{x_1, x_2, x_3\}$
 $m = k = 3$
 $n = 16$



Bloom filter: Lookup

To check membership of an element x in S do:

Lookup(x)


```
for  $i = 1$  to  $k$  do
    compute  $h_i(x)$ 
    if  $B[h_i(x)] == 0$  then
        return no
    end if
end for
return yes
```

Remark 3.

- ▶ If $B[h_i(x)] \neq 1$ for some i , then clearly $x \notin S$.
- ▶ Otherwise, answer “ $x \in S$ ” —*might be a false positive!*

False positive example

Query: “is $x_4 \in S$?”



Lookup(x_4): $h_1(x_4)=h_2(x_4)=h_3(x_4)=1$

Answer: “yes”

Probability of false positive

- ▶ After all elements from S have been hashed into the Bloom filter, the probability that a specific bit is still 0 is

$$\left(1 - \frac{1}{n}\right)^{km} \approx e^{-km/n} = p.$$

- ▶ To simplify the analysis, *assume* that the fraction of bits that are still 0 is **exactly** p .
 - ▶ The fraction of bits is a random variable; we *assume* that it takes a value equal to its expectation.
- ▶ The probability of a false positive is the probability that all k hashes evaluate to 1:

$$f = (1 - p)^k$$

Optimal number of hash functions

$$f = (1 - p)^k = (1 - e^{-km/n})^k$$

- ▶ Trade-off between k and p : using more hash functions
 - ▶ gives us more chances to find a 0 when $x \notin S$;
 - ▶ but reduces the number of 0s in the array!
- ▶ Compute optimal number k^* of hash functions by minimizing f as a function of k :

$$k^* = (n/m) \cdot \ln 2$$

- ▶ Then the **false positive probability** is given by

$$f = (1/2)^{k^*} \approx (0.6185)^{n/m}$$

Big savings in space

- ▶ **Space** required by Bloom filter *per element* of S : n/m bits.
 - ▶ For example, set $n = 8m$. Then $k^* = 6$ and $f \approx 0.02$.
- ⇒ Small constant false positive probability by using only 8 bits (1 byte) per element of S , **independently** of the size of S !

Summary on Bloom filters

Bloom filter can answer approximate set membership in

- ▶ “**constant**” time (time to hash)
- ▶ **constant** space to represent an element from S
- ▶ **constant** false positive probability f .

Application 1 (historical): spell checker

- ▶ Spelling list of 210*KB*, 25*K* words.
- ▶ Use 1 byte per word.
- ▶ Maintain 25*KB* Bloom filter.
- ▶ False positive = accept a misspelled word.

Application 2: implementing joins in database

- ▶ **Join:** Combine two tables with a common domain into a single table.
- ▶ **Semi-join:** A join in distributed DBs in which only the joining attribute from one site is transmitted to the other site and used for selection. The selected records are sent back.
- ▶ **Bloom-join:** A semi-join where we send only a BF of the joining attribute.

Example

| Empl | Sal | Add | City |
|-----------|-----|-----|-----------|
| Bale | 90K | ... | New York |
| Jones | 45K | ... | New York |
| Fletcher | 45K | ... | Pittsburg |
| Rodriguez | 80K | ... | Chicago |
| Shaw | 45K | ... | Chicago |

| City | Cost Of Living |
|-----------|----------------|
| New York | 60K |
| Chicago | 55K |
| Pittsburg | 40K |

Create a table of all employees that make $< 50K$ and live in city where Cost Of Living = COL $> 50K$.

| Empl | Sal | Add | City | COL |
|------|-----|-----|------|-----|
|------|-----|-----|------|-----|

- **Join:** send (City, COL) for COL > 50 .
- **Semi-join:** send just (City) for COL > 50 .
- **Bloom-join:** send a Bloom filter for all cities with COL > 50